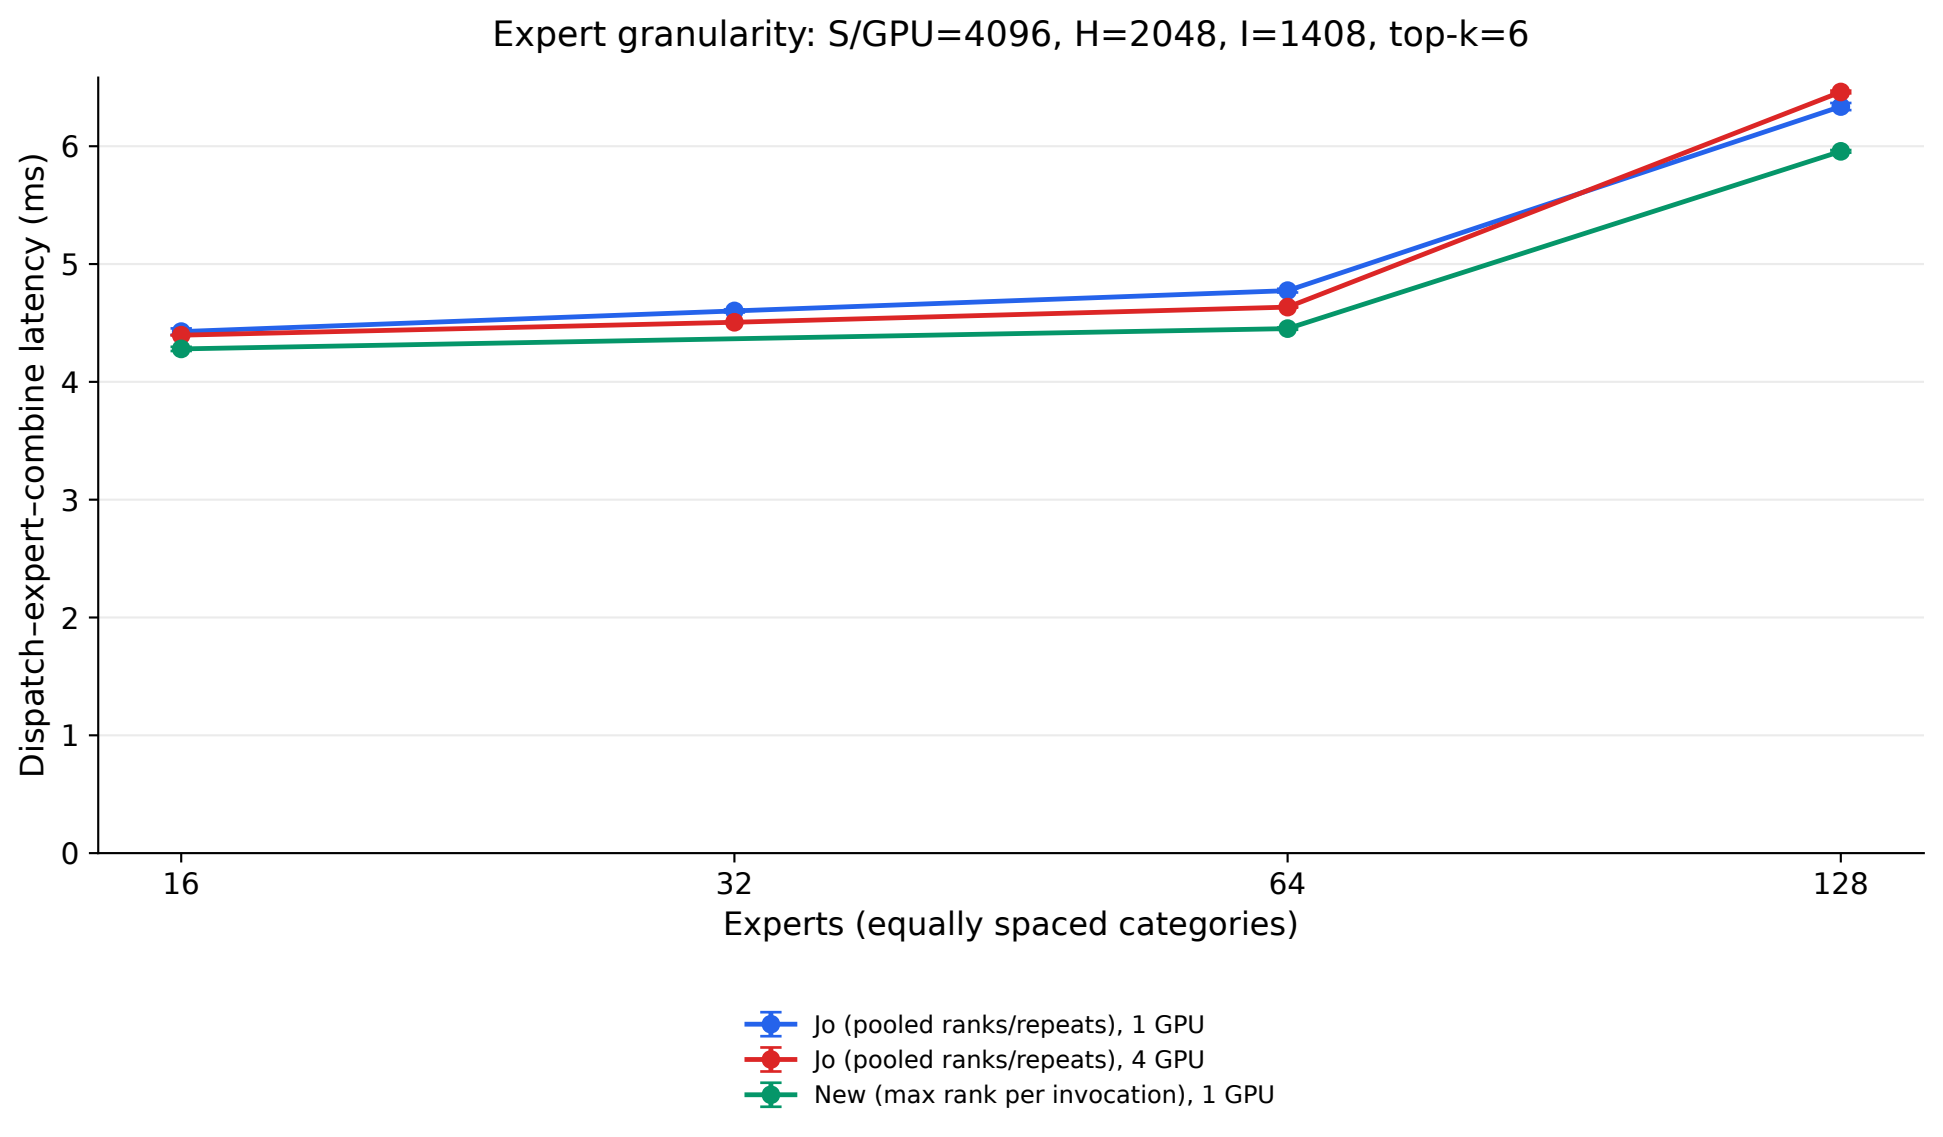
Expert granularity: S/GPU=4096, H=2048, I=1408, top-k=6
Dispatch–expert–combine latency (ms)
0
1
2
3
4
5
6
16
32
64
128
Experts (equally spaced categories)
Jo (pooled ranks/repeats), 1 GPU
Jo (pooled ranks/repeats), 4 GPU
New (max rank per invocation), 1 GPU