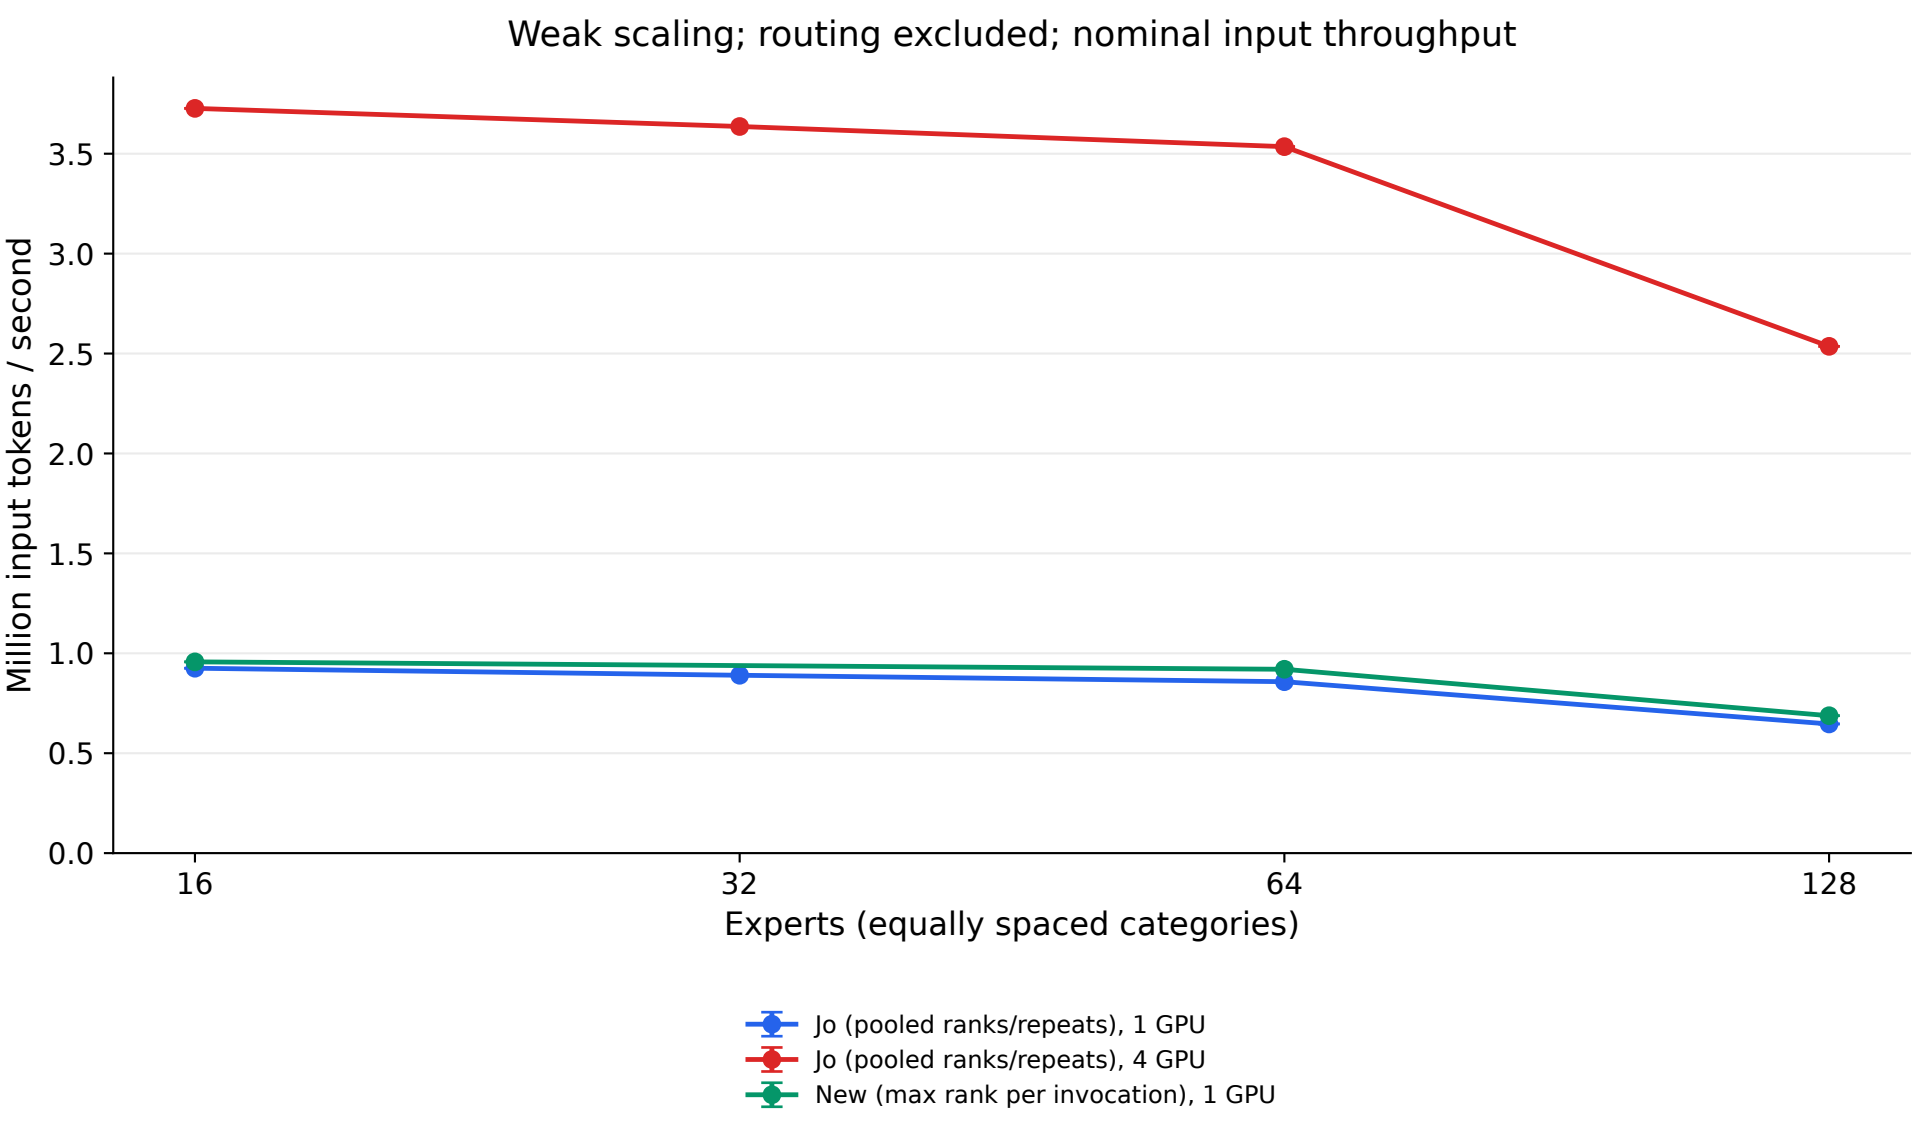
Weak scaling; routing excluded; nominal input throughput
Million input tokens / second
0.0
0.5
1.0
1.5
2.0
2.5
3.0
3.5
16
32
64
128
Experts (equally spaced categories)
Jo (pooled ranks/repeats), 1 GPU
Jo (pooled ranks/repeats), 4 GPU
New (max rank per invocation), 1 GPU