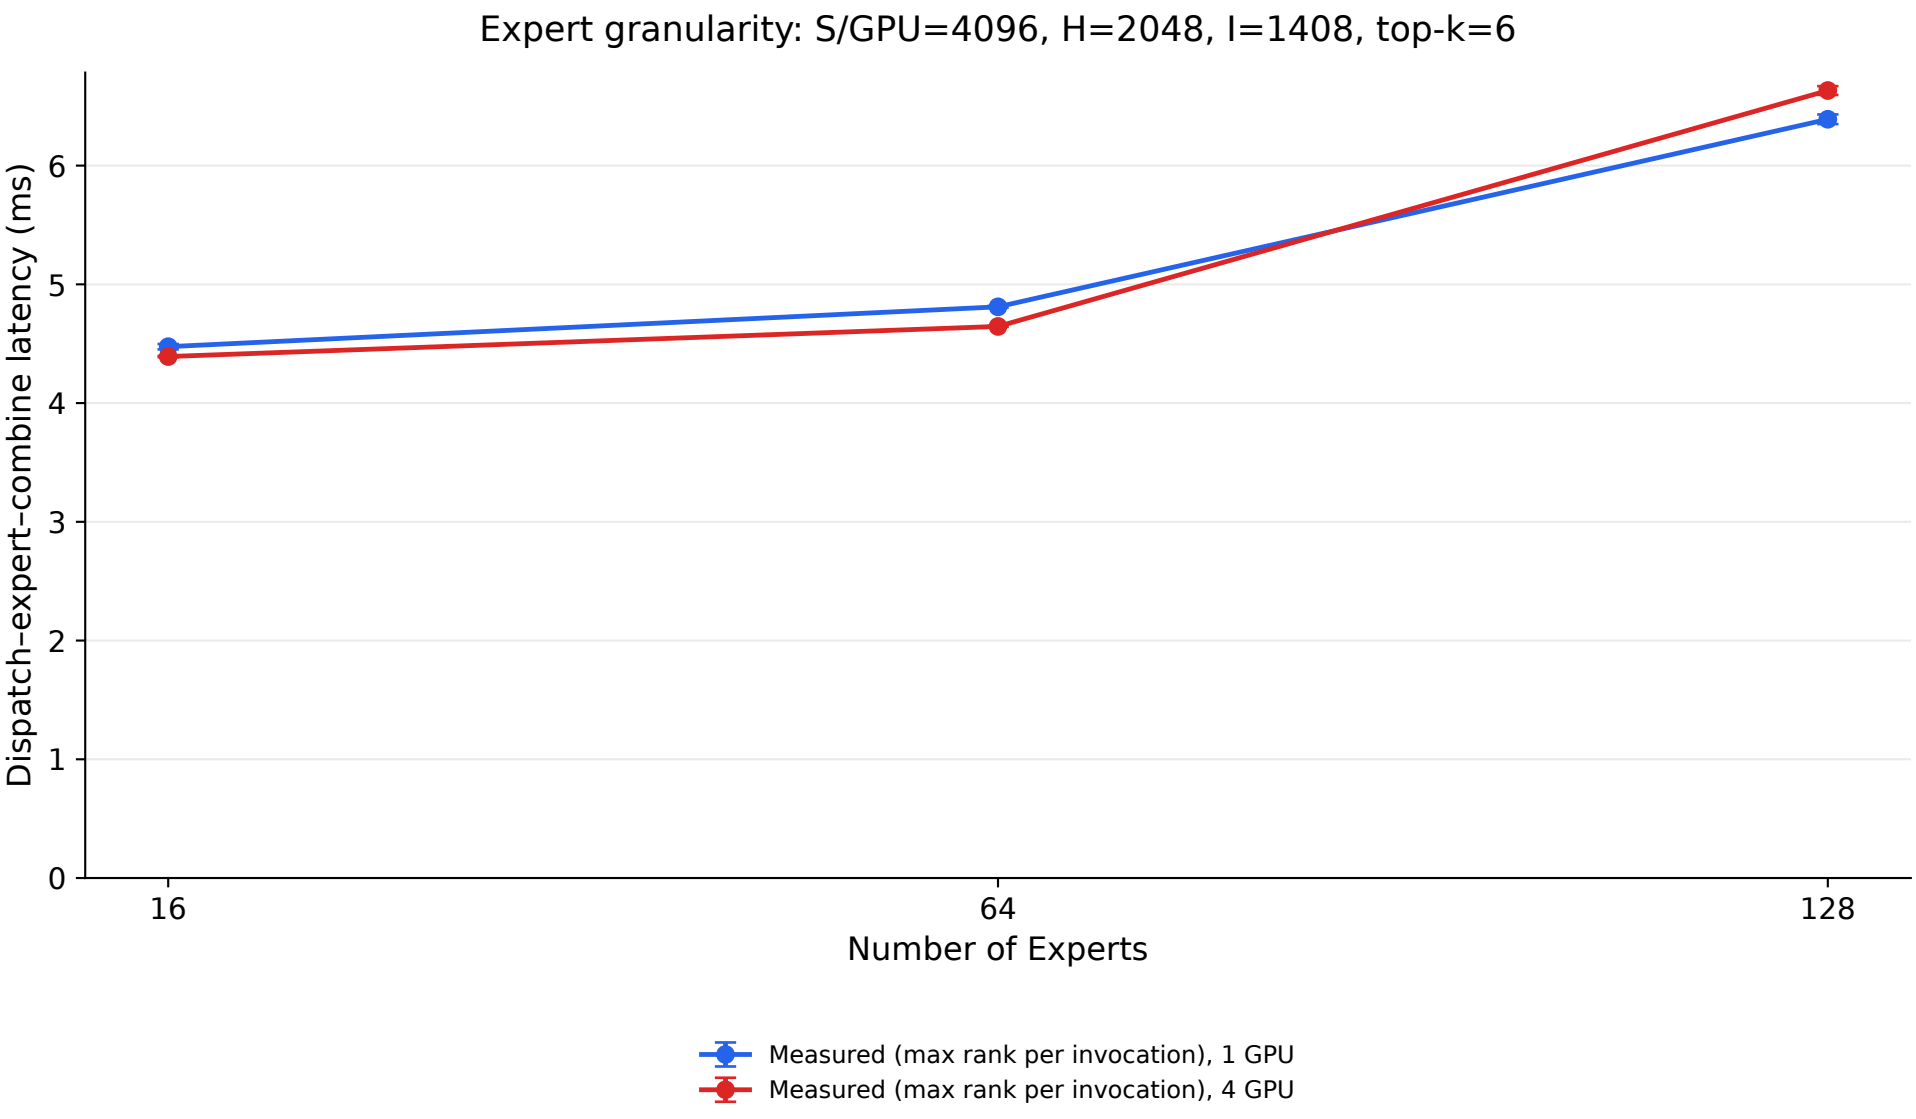
Expert granularity: S/GPU=4096, H=2048, I=1408, top-k=6
Dispatch–expert–combine latency (ms)
0
1
2
3
4
5
6
16
64
128
Number of Experts
Measured (max rank per invocation), 1 GPU
Measured (max rank per invocation), 4 GPU