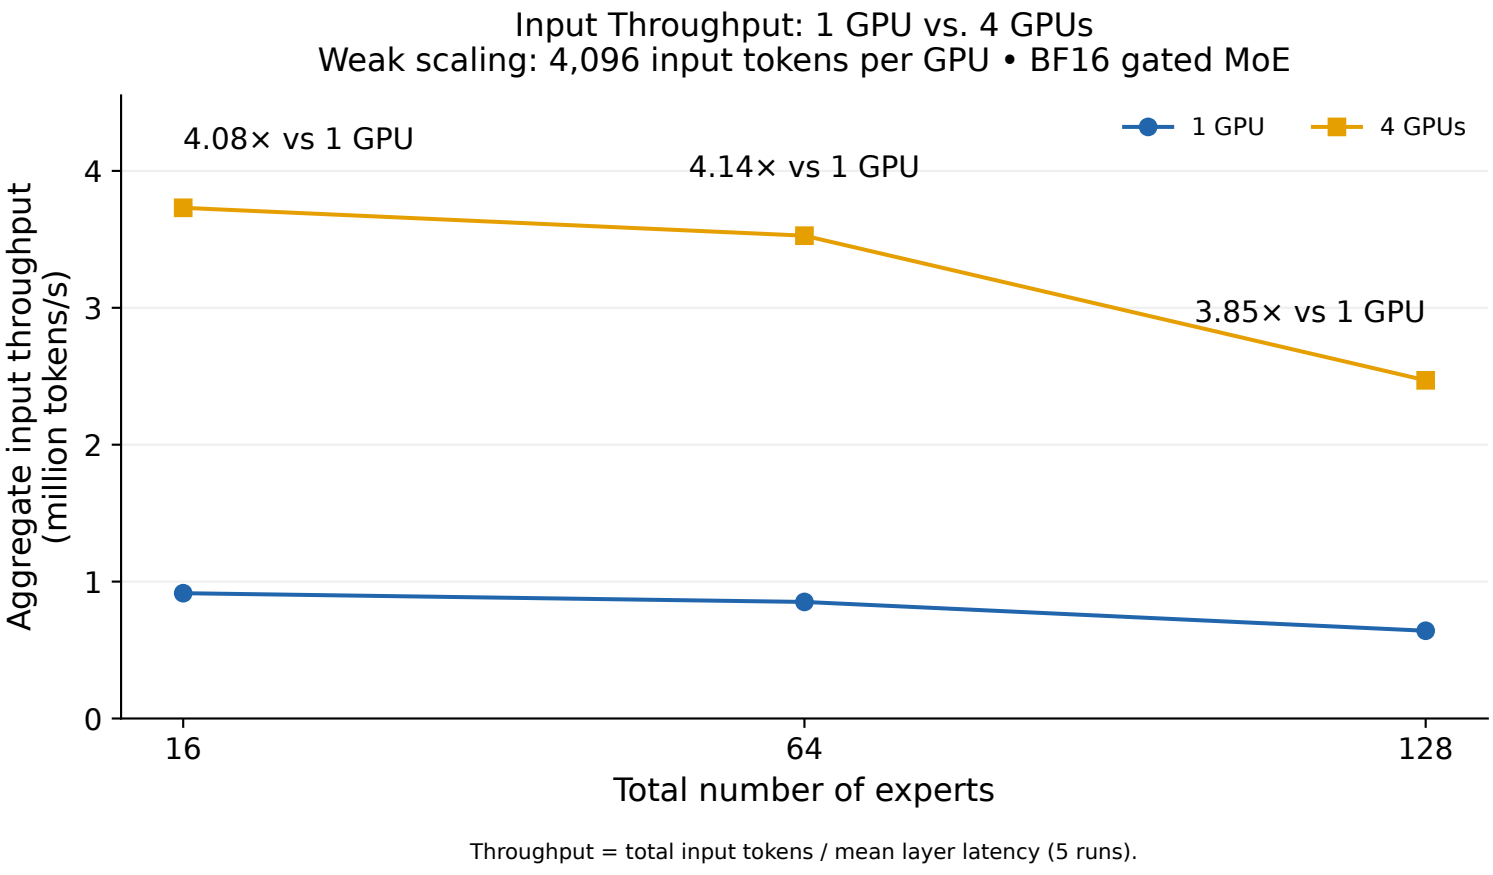
Input Throughput: 1 GPU vs. 4 GPUs
Weak scaling: 4,096 input tokens per GPU • BF16 gated MoE
1 GPU
4 GPUs
4.08× vs 1 GPU
4.14× vs 1 GPU
3.85× vs 1 GPU
Aggregate input throughput
(million tokens/s)
0
1
2
3
4
16
64
128
Total number of experts
Throughput = total input tokens / mean layer latency (5 runs).