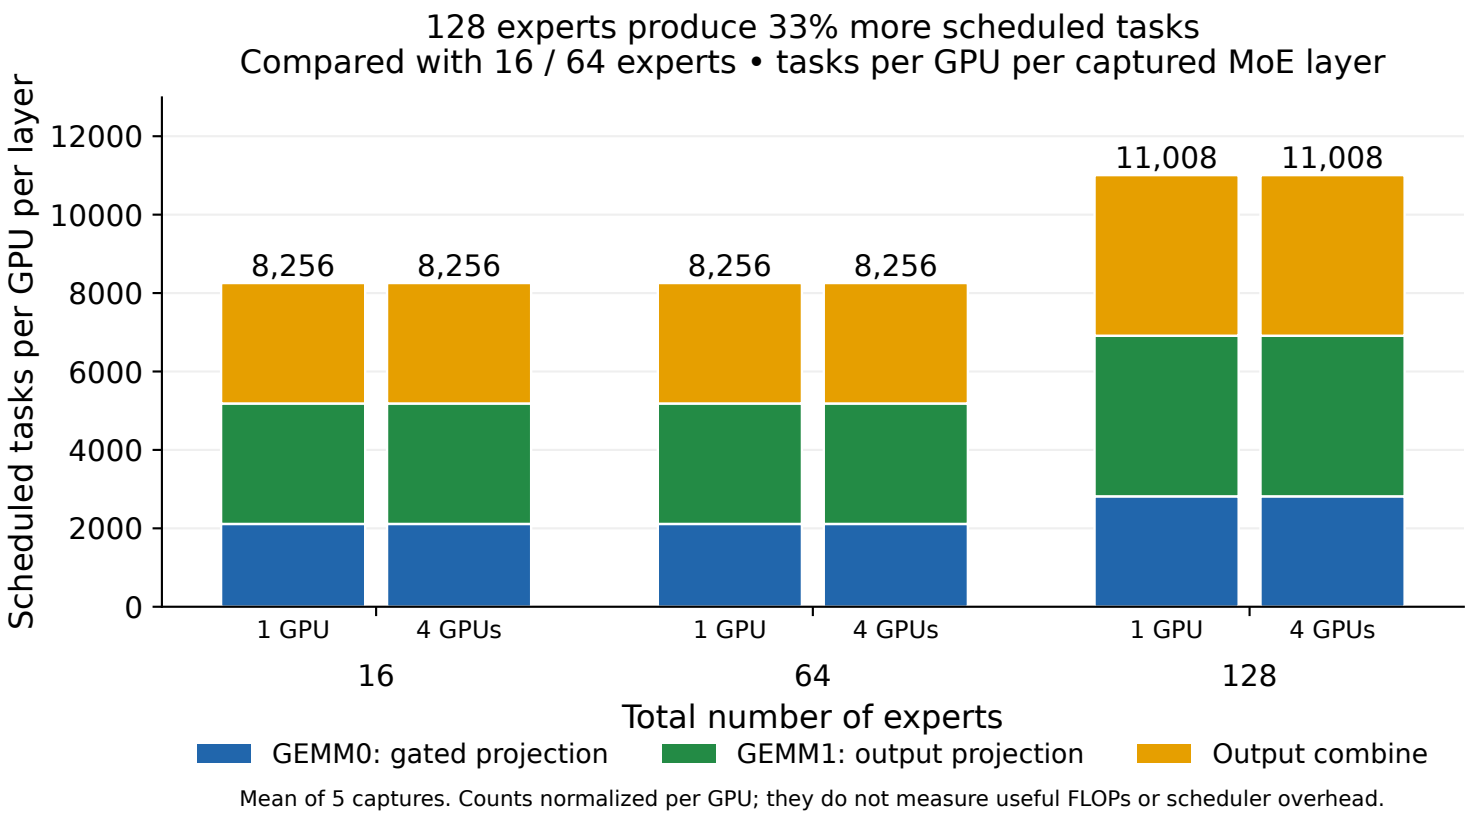
128 experts produce 33% more scheduled tasks
Compared with 16 / 64 experts • tasks per GPU per captured MoE layer
Scheduled tasks per GPU per layer
0
2000
4000
6000
8000
10000
12000
8,256
8,256
8,256
8,256
11,008
11,008
1 GPU
4 GPUs
1 GPU
4 GPUs
1 GPU
4 GPUs
16
64
128
Total number of experts
GEMM0: gated projection
GEMM1: output projection
Output combine
Mean of 5 captures. Counts normalized per GPU; they do not measure useful FLOPs or scheduler overhead.